

Deep Learning

MSDS 631

Images and
Convolutional Neural Networks

Michael Ruddy

Questions?

- From last lecture?
- From the lab assignment?

Overview

- Why imaging? Imaging tasks?
- What/Why is a Convolution?
- Why text? Text tasks?
- Preprocessing Text

Unstructured Data



- Often refer to tabular data as “structured data”
- Thus, we refer to images and text as *unstructured data*
 - Also include things like video (lots of images) and audio (fancy sequence)
- Structured Data
 - Random Forest vs. Gradient Boosting vs. DL vs. etc
- Unstructured Data
 - Deep Learning reigns supreme

Why Imaging?

- Humans are really good at looking at things
 - The human eye/brain is an incredibly complicated piece of machinery
- Efforts to recreate vision based on human models of vision have largely been unsuccessful

Why Imaging?

- Humans are really good at looking at things
 - The human eye/brain is an incredibly complicated piece of machinery
- Efforts to recreate vision based on human models of vision have largely been unsuccessful
- Imaging is important!



Why Imaging?

- Humans are really good at looking at things
 - The human eye/brain is an incredibly complicated piece of machinery
- Efforts to recreate vision based on human models of vision have largely been unsuccessful
- Imaging is important!
 - Classification (facial/object recognition, avoid poisonous plants, etc.)
 - Medical Imaging (detecting disease, predicting outcomes of radiation, segmentation of medical images)
 - Autonomous Driving (driver assistance, fully autonomous vehicles)
 - Deepfakes and deepfake detection

Why Imaging?

- Humans are really good at looking at things
 - The human eye/brain is an incredibly complicated piece of machinery
- Efforts to recreate vision based on human models of vision have largely been unsuccessful
- Imaging is important!
 - Classification (facial/object recognition, avoid poisonous plants, etc.)
 - Medical Imaging (detecting disease, predicting outcomes of radiation, segmentation of medical images)
 - Autonomous Driving (driver assistance, fully autonomous vehicles)
 - Deepfakes and deepfake detection
- A lot of these are time-consuming things that human can do really well

Why are images special?

- Images are deceptively hard



Why are images special?

- Images are deceptively hard

This is a cat



Why are images special?

- Images are deceptively hard

This is ??????

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ .5 & .75 & 1 & \dots & .25 \\ \vdots & \vdots & \vdots & & \vdots \\ .333 & 0 & 1 & \dots & 0 \end{bmatrix}$$

Why are images special?

- Images are deceptively hard
- Images are big



32x32 image
1024 features



512x512 image
262,144 features

Why are images special?

- Images are deceptively hard
- Images are big



32x32 image
1024 features



512x512 image
262,144 features

Fully Connected Layer

- $1024 \rightarrow 1024$
- $1024^2 = 1,048,576$
parameters
- $262,144 \rightarrow 262,144$
- $68,719,476,736$
parameters

Why are images special?

- Images are deceptively hard
- Images are big
- Geometry matters!
 - Pixels near each other interact in different ways to create features than pixels far away

Different
relationships

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ .5 & .75 & 1 & \dots & .25 \\ \vdots & \vdots & \vdots & & \vdots \\ .333 & 0 & 1 & \dots & 0 \end{bmatrix}$$

Why are images special?

- Images are deceptively hard
- Images are big
- Geometry matters!
 - Pixels near each other interact in different ways to create features than pixels far away
 - This is free data that we lose if we simply consider an image as a data vector

Different
relationships

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ .5 & .75 & 1 & \dots & .25 \\ \vdots & \vdots & \vdots & & \vdots \\ .333 & 0 & 1 & \dots & 0 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data

The Convolution

- Fancy **linear** operation useful for spatial data

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

The Convolution

- Fancy **linear** operation useful for spatial data

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} & & & \\ & & & \\ & & & \\ & & & \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(1 \times 1) + (.5 \times 0) + (0 \times 0) + (.25 \times 2) \\ = 1.5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

=

$$\begin{bmatrix} 1.5 & \dots & \dots \\ \vdots & \ddots & \vdots \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(.5 \times 1) + (1 \times 0) + (.25 \times 0) + (.5 \times 2) = 1.5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

$$= \begin{bmatrix} 1.5 & 1.5 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(1 \times 1) + (0 \times 0) + (.5 \times 0) + (1 \times 2) = 3$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*

$$= \begin{bmatrix} 1.5 & 1.5 & \rightarrow 3 \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

$$(0 \times 1) + (.25 \times 0) + (1 \times 0) + (.25 \times 2) = .5$$

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

*



$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & & \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & ? & ? \\ ? & ? & ? \end{bmatrix}$$

The Convolution

- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

Filter
(kernel)

*

$$\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

=

$$\begin{bmatrix} 1.5 & 1.5 & 3 \\ 0.5 & .25 & 2.5 \\ 1 & 2.25 & 2 \end{bmatrix}$$

The Convolution


- Fancy **linear** operation useful for spatial data
- Element-wise product

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{[Output Image]}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?

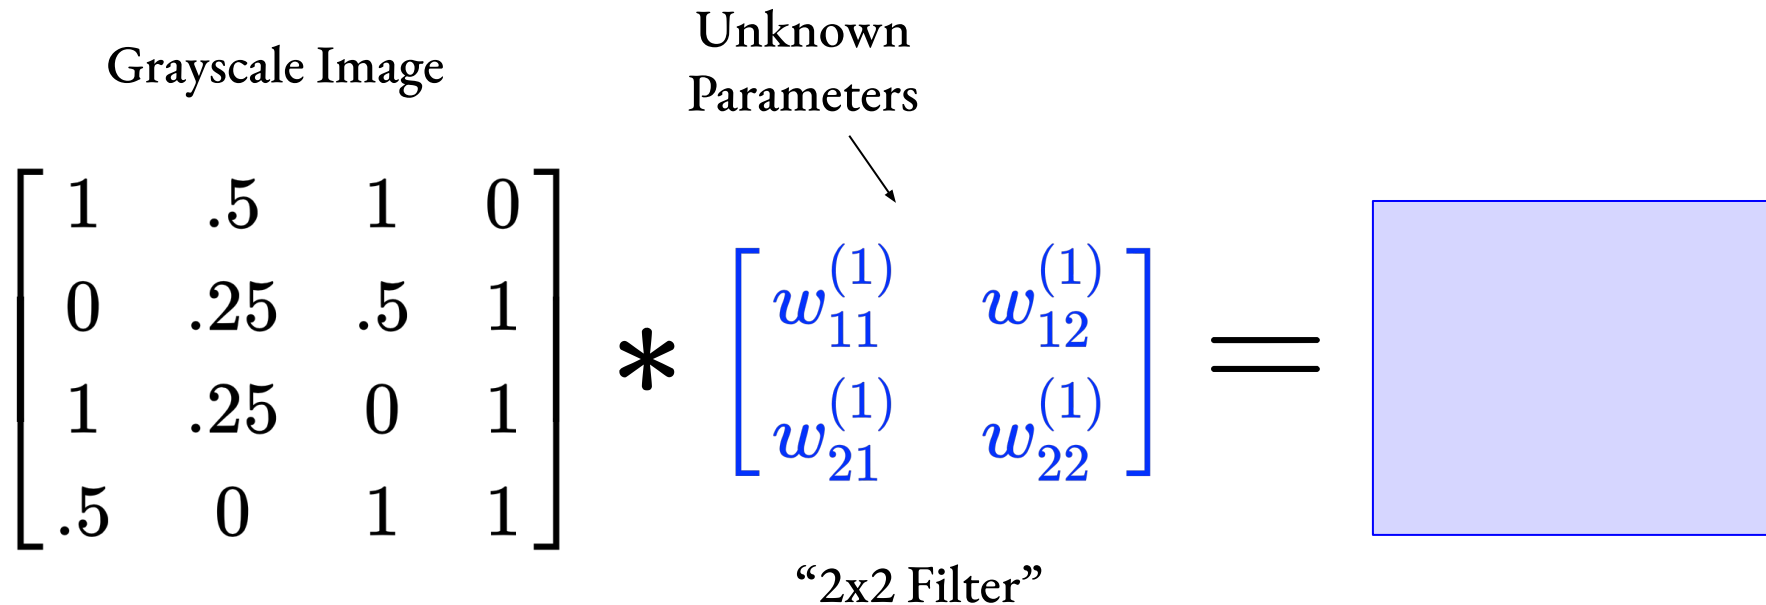
- Only four parameters!
 - If input is dimension 16 and output is dimension 9, how many for FC?

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{Output Image}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?


- Only four parameters!
- Translational Equivariance
 - If I shift my image, I shift the output!

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{Image}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?


- Only four parameters!
- Translational Equivariance
- Weight Sharing (detect same feature translated to different parts of the image)

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{[Output Box]}$$

Unknown Parameters

“2x2 Filter”



Why Convolution?

Intuition: Edge
Detection

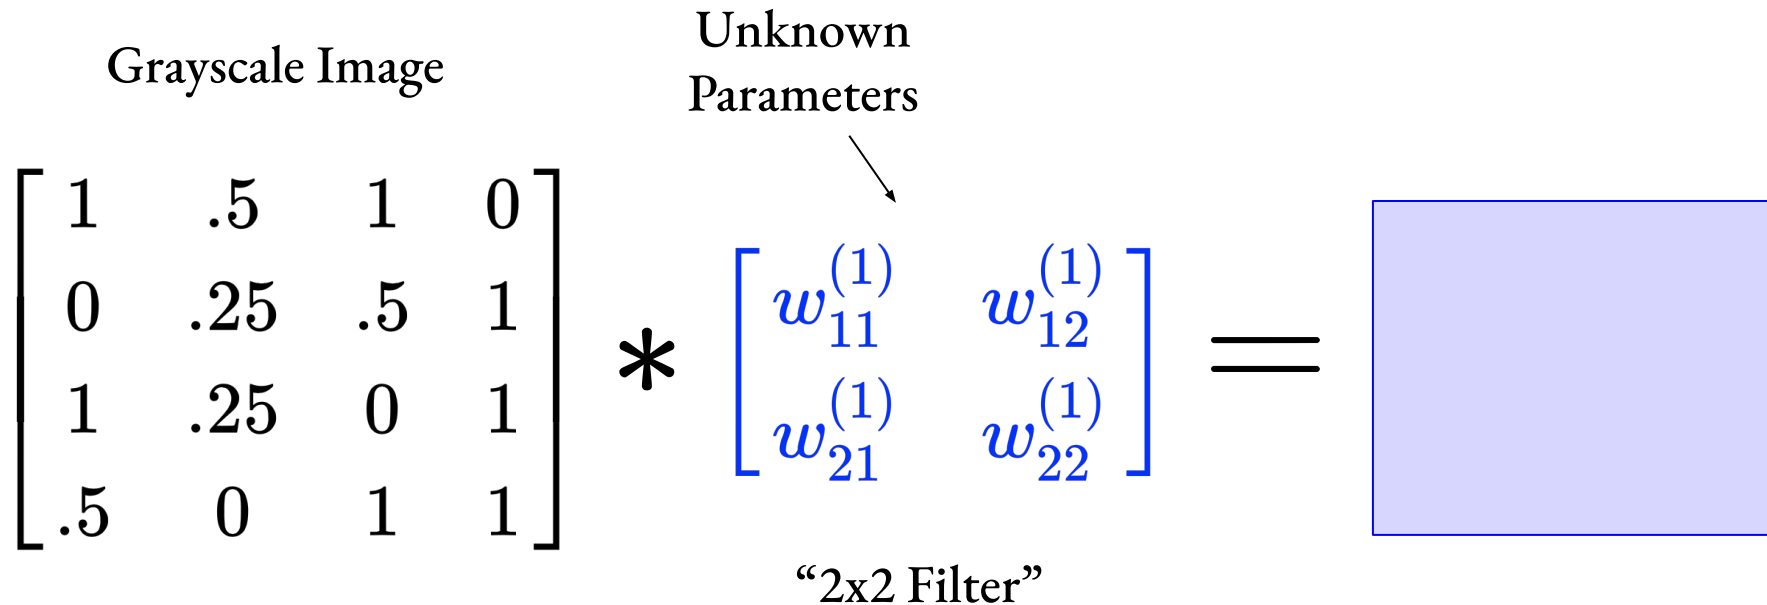
- Only four parameters!
- Translational Equivariance
- Weight Sharing (detect same feature translated to different parts of the image)

Grayscale Image

$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} \end{bmatrix} = \text{[Output Box]}$$

Unknown Parameters

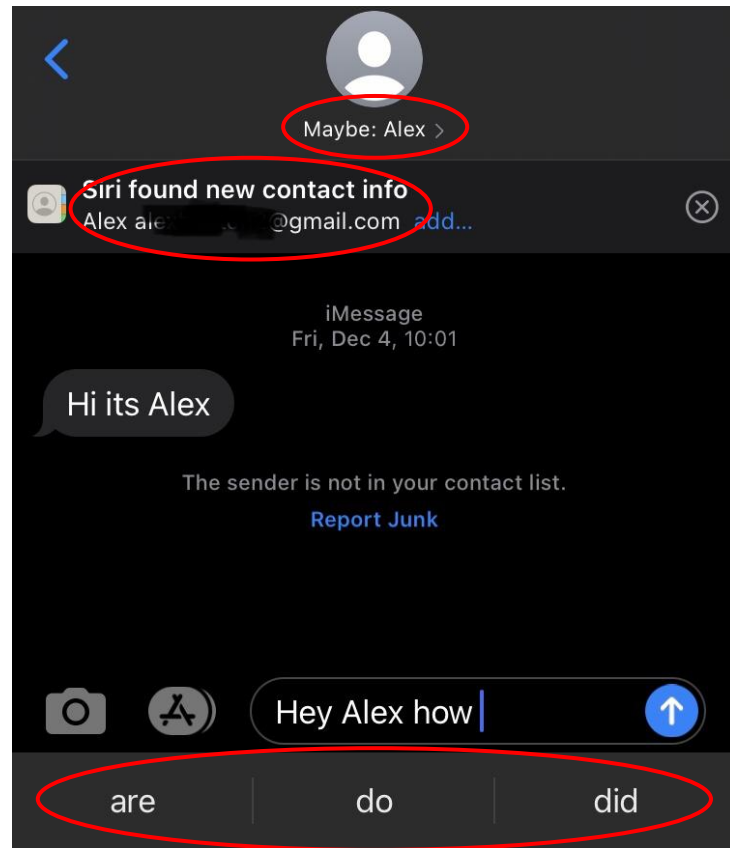
“2x2 Filter”

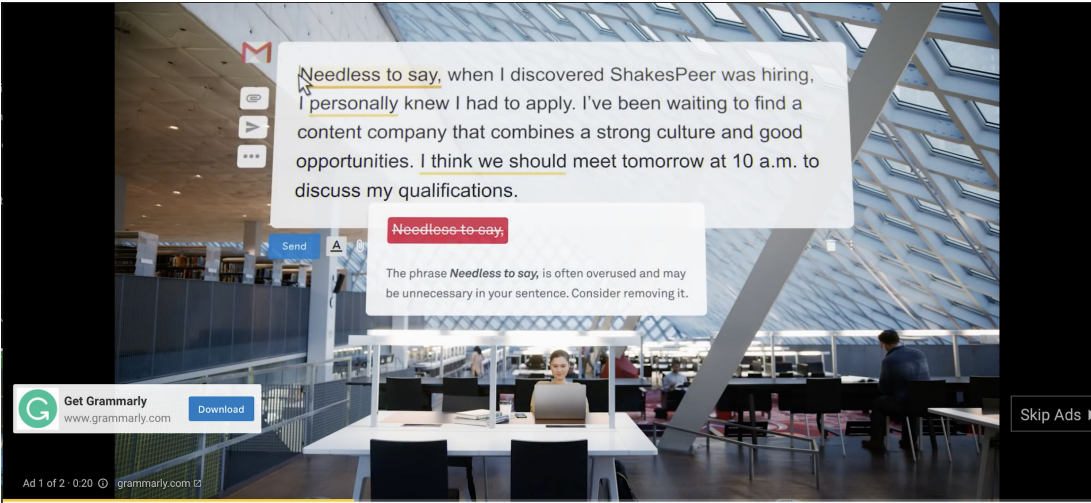


Why Natural Language Processing?

- Understand, analyze, and perform tasks using human language (through text).
- Example Tasks:
 - Sentiment Analysis
 - Auto-complete
 - Translation
 - Question answering
 - Conversation?!

Some or all of the content shared in this Tweet conflicts with guidance from public health experts regarding COVID-19. [View](#) [Learn more](#)





Needless to say, when I discovered ShakesPeer was hiring, I personally knew I had to apply. I've been waiting to find a content company that combines a strong culture and good opportunities. I think we should meet tomorrow at 10 a.m. to discuss my qualifications.

Needless to say,

The phrase **Needless to say**, is often overused and may be unnecessary in your sentence. Consider removing it.

Get Grammarly
www.grammarly.com

Download

Send

Skip Ads ▶

Ad 1 of 2 · 0:20 · grammarly.com

How to write descriptively - Nalo Hopkinson

3,057,044 views · Nov 16, 2015

👍 105K 👎 1.3K ➦ SHARE ⚙️ SAVE ...

Find the words. ✓

English ↔ German

I love and hate learning German ✕

Ich liebe und hasse es, Deutsch zu lernen ➡

NLP is hard!

- How to represent text as data?



$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

NLP is hard!

- How to represent text as data?
- Humans represent text using characters
 - Takes years to learn to read
 - Different peoples do it differently all around the world



$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

train

brain

head

NLP is hard!

- How to represent text as data?
- Humans represent text using characters
 - Takes years to learn to read
 - Different peoples do it differently all around the world



$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

train \longrightarrow 20-18-1-9-14

brain \longrightarrow 2-18-1-9-14

head \longrightarrow 8-5-1-4

NLP is hard!

- How to represent text as data?
- Humans represent text using characters
 - Takes years to learn to read
 - Different peoples do it differently all around the world
- For most tasks this is not a particularly helpful embedding
 - Intrinsic meaning is largely lost



$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$

train → 20-18-1-9-14

brain → 2-18-1-9-14


head → 8-5-1-4

NLP is hard!

- How to represent text as data?
- Humans represent text using characters
 - Takes years to learn to read
 - Different peoples do it differently all around the world
- For most tasks this is not a particularly helpful embedding
 - Intrinsic meaning is largely lost



$$\begin{bmatrix} 1 & .5 & 1 & 0 \\ 0 & .25 & .5 & 1 \\ 1 & .25 & 0 & 1 \\ .5 & 0 & 1 & 1 \end{bmatrix}$$



train → 20-18-1-9-14

brain → 2-18-1-9-14

head → 8-5-1-4

Tokenization

- Idea: Break up text into pieces (tokens) and treat as categorical variables
 - Often these tokens are words

Tokenization


- Idea: Break up text into pieces (tokens) and treat as categorical variables
 - Often these tokens are words

$$\begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} \quad \begin{bmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix} \quad \cdots \quad \begin{bmatrix} 0 \\ 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}$$


a

aadvark



zebra



Word Embedding



High-dimensional space



Low-dimensional space

Other Types of Tokenization

- Idea: Break up text into pieces (tokens) and treat as categorical variables
- Words -> Tokens

Other Types of Tokenization

- Idea: Break up text into pieces (tokens) and treat as categorical variables
- Words -> Tokens
 - N-grams: Common phrases as one token instead of separate tokens


data_science


vs.


data, science

Other Types of Tokenization

- Idea: Break up text into pieces (tokens) and treat as categorical variables
- Words -> Tokens
- Characters -> Tokens

train  20-18-1-9-14

brain  2-18-1-9-14

head  8-5-1-4

Other Types of Tokenization

- Idea: Break up text into pieces (tokens) and treat as categorical variables
- Words -> Tokens
- Characters -> Tokens
- Sub-words -> Tokens
 - Break up words into smaller tokens
 - Smaller dictionary, less total tokens
 - Better at handling unknown, less lemmatization

Unfortunately -> un + fortunate + ly
skiing -> ski + ing

Other Types of Tokenization

- Idea: Break up text into pieces (tokens) and treat as categorical variables
- Words -> Tokens
- Characters -> Tokens
- Sub-words -> Tokens
 - Break up words into smaller tokens
 - Smaller dictionary, less total tokens
 - Better at handling unknown, less lemmatization
 - Many Algorithms: BPE, Unigram, WordPiece

Unfortunately -> un + fortunate + ly
skiing -> ski + ing

Other Types of Tokenization

- Idea: Break up text into pieces (tokens) and treat as categorical variables
- Words -> Tokens
- Characters -> Tokens
- Sub-words -> Tokens
- Sentence Segmentation
 - EOS (End of Sentence) and SOS (Start of Sentence) tokens are common

Other Types of Tokenization

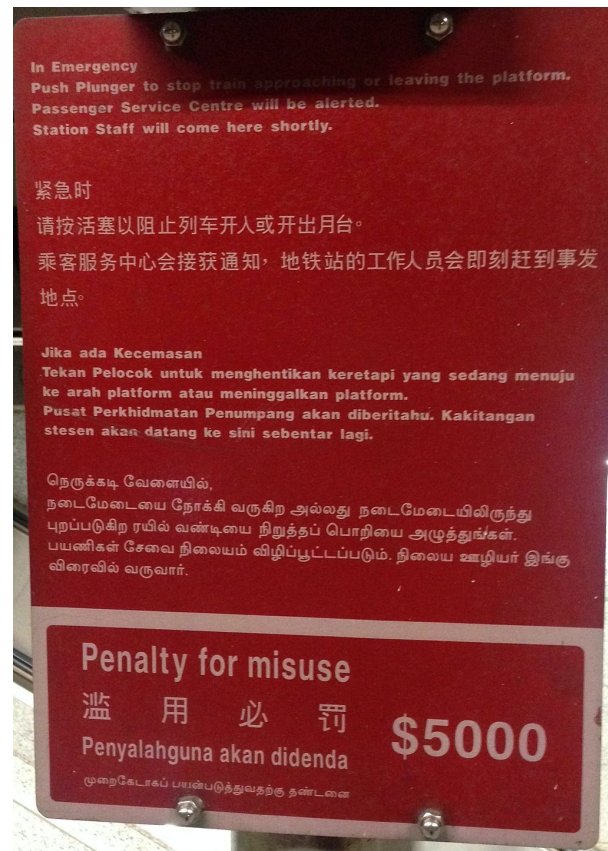
- Idea: Break up text into pieces (tokens) and treat as categorical variables
- Words -> Tokens
- Characters -> Tokens
- Sub-words -> Tokens
- Sentence Segmentation
 - EOS (End of Sentence) and SOS (Start of Sentence) tokens are common
 - Non-trivial to find these!
 - Binary Classifier, complicated logic trees

Can't just
rely on
periods!

The U.K. exports of goods and services as percent of GDP was 31.6% in 2019.

Other Types of Tokenization

- Idea: Break up text into pieces (tokens) and treat as categorical variables
- Words -> Tokens
- Characters -> Tokens
- Sub-words -> Tokens
- Sentence Segmentation
- Other languages:
 - Chinese languages, Arabic, French, etc.



Word Tokenization Techniques

- Lemmatization
 - Reduce words to their base
 - Shrink dictionary size

running -> run
mice -> mouse

Word Tokenization Techniques

- Lemmatization
- Infrequent words (misspelled or weird words)
 - Remove from text or encode as single UNK token

Word Tokenization Techniques

- Lemmatization
- Infrequent words (misspelled or weird words)
- Cleaning before tokenization
 - Lower case
 - Remove weird characters/numbers/punctuation
 - Remove stop words

the, to, a, an, etc.

Word Tokenization Techniques

- Lemmatization
- Infrequent words (misspelled or weird words)
- Cleaning before tokenization
 - Lower case
 - Remove weird characters/numbers/punctuation
 - Remove stop words
- Named Entity Recognition



Apple vs. apple
Xerox vs. xerox



Deep Learning and NLP

- Sequences
 - Variable length
 - Relationships between elements of sequence
- Continuous Bag of Words (CBOW)
- 1D CNN
- Recurrent Neural Network (RNN)
 - Keep track of a hidden state vector of features as you move along a sequence
 - Sequence length agnostic

Summary

- Images and Text are special
 - Humans are better at seeing than speaking
 - Language is “harder” than vision
- Special architectures exist to take advantage of the unique properties
 - Images: spatial
 - Text: sequences
- NLP requires a lot of preprocessing and thinking deeply about representation